

Communication-Aware Policy Optimization with Adaptive Topology Control for Multi-UAV Cooperative Navigation

Abstract—Multi-UAV systems require efficient cooperation under limited communication resources, including bandwidth, energy, and latency constraints. Existing multi-agent reinforcement learning methods typically assume dense or geometry-predefined messaging and therefore do not treat communication itself as an optimization target. We propose AC-DSGF, which learns a task-dependent communication topology jointly with residual guidance under an explicit communication cost and Top- K budget. On 16-UAV navigation with five seeds, AC-DSGF achieves comparable task performance to DSGF while reducing communication overhead by over one order of magnitude and improving the Communication Efficiency Index. Additional budget, silence-collapse, trigger, and packet-loss studies indicate task-driven sparsity and graceful degradation under unreliable links, supporting bandwidth-constrained cooperative UAV deployment.

Index Terms—Multi-UAV systems, multi-agent reinforcement learning, communication-efficient cooperation, dynamic graphs, residual guidance.

I. INTRODUCTION

a) Engineering reality.: In multi-UAV cooperation, every message costs bandwidth, energy, and latency. Controllers that optimize only local policies under an implicit “unlimited radio” assumption therefore leave a first-class resource unoptimized: **communication itself**.

b) Limitation of prior MARL guidance.: MAPPO-style CTDE stabilizes decentralized control [1], and graph guides (e.g., GAT) encode spatial relations for early coordination. A common pipeline is

policy $\rightarrow$ (predefined) communication $\rightarrow$ environment.

Topology is treated as given—by a radius, a dense graph, or a hand schedule—so agents cannot decide *whether* an exchange is currently needed.

c) Our shift.: We invert that dependency and learn topology as part of decision making:

environment $\rightarrow$ learn communication topology $\rightarrow$ policy.

Unlike previous approaches that assume communication topology is predefined, we treat topology selection as a learnable decision process.

Concretely, AC-DSGF answers: *when* to talk, *with whom*, and *how much* under a budget (Q1–Q3).

d) Method sketch.: **AC-DSGF** extends DSGF v2 with a soft gate g_{ij}^t , a Top- K budget layer $A_t^{\text{AC}} = A \odot E^K$, and the joint objective $\mathbb{E}[\sum r_t] - \lambda_c \mathbb{E}[\sum C_t]$, while residual guidance $a = \pi(o) + \beta \Delta(\Phi)$ keeps communication assistive rather than dominant (Fig. 1).




Fig. 1. Motivation: without residual decoupling, guided MARL can collapse under long-horizon optimization (4-UAV ablation; Table VI). AC-DSGF keeps residual guidance while adding budget-aware communication.

e) Claim.: We do not position AC-DSGF as a Success maximizer. Experiments show it **achieves comparable task performance** to DSGF while **reducing communication overhead by over one order of magnitude** (Table V).

f) Contributions.:

- 1) **Learnable communication topology**: a soft gate $g_{ij}^t = \sigma(f_\theta(h_i, h_j, d_{ij}, q_{ij}))$ on radius support.
- 2) **Budget-aware edge selection**: Top- K adjacency $A_t^{\text{AC}} = A \odot E_t^K$ under $J = \mathbb{E}[\sum r] - \lambda_c \mathbb{E}[\sum C]$.
- 3) **Residual policy decoupling**: $a = \pi(o) + \beta \Delta(\Phi)$ so guidance assists rather than dominates execution.
- 4) **Evaluation** at $N=16$ (five seeds) covering comparison, budget, silence-collapse, trigger/stability, compute cost, and packet-loss robustness.

II. RELATED WORK

We organize prior art into three lines and highlight the residual gap that AC-DSGF targets: *joint* optimization of task return, adaptive topology, and explicit communication budgets.

A. Multi-Agent RL for UAV Swarms

Centralized Training with Decentralized Execution (CTDE) methods such as MAPPO stabilize cooperative control by sharing a critic during training while keeping actors local at execution time [1], [2], [3]. Extensions to aerial swarms demonstrate strong local navigation and collision handling under partial observability [4], [5]. However, most CTDE pipelines optimize policy return under an implicit assumption of **sufficient communication resources**: either agents are treated as independent given a shared training signal, or coordination relies on unlimited centralized features that do

not survive bandwidth, latency, and energy limits on real radio links.

Gap 1. Existing MARL methods for UAV cooperation mainly optimize policy performance while assuming sufficient communication resources.

B. Graph-based Multi-Agent Interaction Modeling

Graph neural networks and attention mechanisms inject relational inductive bias into MARL [6], [7], [8], [9]. GAT-style guides and graph Transformers adaptively weight neighbors and can capture time-varying spatial relations for swarm navigation. Yet attention support is typically defined by geometry (e.g., communication radius) or is fully dense. Dense message passing incurs $\mathcal{O}(N^2)$ cost that grows quickly for $N \geq 16$ –32 UAVs, while binary radius graphs still exchange messages on *all* feasible edges regardless of task urgency. Our prior DSGF improves sparsity via quality weighting and residual guidance, but still activates the full geometric support whenever $A_{ij}=1$.

Gap 2. Graph MARL models interaction topology but lacks **explicit communication budget optimization**—topology adaptation is not equated with bandwidth-aware economy.

C. Communication-Efficient MARL

A parallel literature studies message pruning, gated communication, dropout, and bandwidth-aware objectives to reduce exchange volume [10], [11], [12], [13], [14], [15], [16]. Many approaches, however, rely on random or heuristic suppression of messages, hand-designed schedules, or discrete “talk/silence” decisions that are weakly coupled to continuous residual guidance for control. Under severe budgets, random pruning can discard coordination-critical links; conversely, over-aggressive silence collapses information flow.

Gap 3. Prior communication-efficient MARL often **drops messages** rather than *learning when communication is necessary* for the ongoing task under a quantified budget.

a) Positioning. Different from the above, AC-DSGF jointly considers (i) task performance, (ii) *learned* communication topology adaptation via gates g_{ij}^t , and (iii) strict communication-budget constraints through $C_{\text{comm}} = \sum g_{ij}$ and optional Top- K scheduling, while retaining residual DSGF guidance for long-horizon stability. It is therefore not a mere DSGF model with an added L2-style penalty, but a communication-constrained coordination stack whose sparsity is task-dependent (validated in Secs. VI-D–VI-E).

III. PROBLEM FORMULATION

A. Multi-UAV Cooperative Navigation

Consider a team of N UAVs performing cooperative navigation in a continuous 2D workspace under partial observability. At discrete time t , agent $i \in \{1, \dots, N\}$ receives a local observation

$$o_i^t = [p_i^t, v_i^t, g_i^t, \ell_i^t] \in \mathbb{R}^{d_o}, \quad (1)$$

where $p_i^t, v_i^t \in \mathbb{R}^2$ are position and velocity, g_i^t is the goal-relative vector, and ℓ_i^t denotes onboard range sensing. Each agent executes a continuous action $a_i^t \in \mathcal{A}$ under a decentralized policy with centralized training (CTDE).

The team seeks to maximize expected cumulative reward under communication cost,

$$\max_{\theta} \mathbb{E} \left[\sum_{t=0}^H r_t \right] - \lambda_c \mathbb{E} \left[\sum_{t=0}^H C_t \right], \quad r_t = \sum_{i=1}^N r_i^t, \quad (2)$$

where $r_i^t = r_i^{\text{goal}}(t) + r_i^{\text{coll}}(t)$ and C_t is defined below. When $\lambda_c=0$, Eq. (2) reduces to standard return maximization (baselines). For AC-DSGF, $\lambda_c > 0$ makes communication a first-class training signal.

B. Communication Constraint

Airborne links are range-limited and capacity-constrained. Two agents may exchange messages only if $\|p_i^t - p_j^t\| < R_c$, inducing a feasible mask $A^t \in \{0, 1\}^{N \times N}$. We further impose a soft communication budget through a learnable gate $g_{ij}^t \in [0, 1]$. Unless stated otherwise, the communication cost used in tables is the *time-averaged soft gate mass*

$$C_t = \sum_{i \neq j} g_{ij}^t, \quad C = C_{\text{comm}} = \frac{1}{T} \sum_{t=1}^T C_t = \frac{1}{T} \sum_{t=1}^T \sum_{i \neq j} g_{ij}^t. \quad (3)$$

For un-gated baselines (GAT/DSGF), we report the analogous average number of active geometric edges $\frac{1}{T} \sum_t \sum_{i,j} A_{ij}^t$ (or quality-weighted support) so that Comm columns are comparable as *message intensity*, not raw radio packet counts. The Table I AC-DSGF value $C \approx 0.43$ is therefore the mean soft mass under Eq. (3) over the 200-ep evaluation—not an undefined unit. Eq. (2) pairs this C with task return via λ_c . Hard budgets additionally restrict support via $A_t^{\text{AC}} = A^t \odot E_t^K$ (Sec. IV-B).

C. Metric: Communication Efficiency Index

To compare methods under different messaging intensities, we report

$$\text{CEI} = \frac{S}{C_{\text{comm}} + \varepsilon}, \quad \varepsilon = 10^{-6}, \quad (4)$$

where S is the evaluation success rate. MAPPO has no explicit graph messaging ($C_{\text{comm}}=0$) and CEI is marked as not applicable. Our primary claim is *not* that AC-DSGF maximizes S unconditionally, but that it preserves **comparable** S while substantially reducing C_{comm} relative to fixed-topology guidance.

NOTATION


IV. PROPOSED METHOD: AC-DSGF

We propose *Adaptive Communication-aware Dynamic Sparse Graph Fusion* (AC-DSGF). Unlike post-hoc message pruning, AC-DSGF **jointly** learns the control policy and a task-dependent communication topology under an explicit cost. Fig. 2 shows three layers: (i) observation embedding, (ii) adaptive communication (gate + budget), (iii) residual DSGF policy.

TABLE I
SUMMARY OF KEY SYMBOLS. SOFT GATES $g_{ij}^t \in [0, 1]$ ARE CONTINUOUS (PROBABILITY-LIKE INTENSITIES), NOT HARD BINARY EDGES; BINARY DECISIONS APPEAR ONLY AFTER TOP- K WHEN A HARD BUDGET IS ENFORCED.

Symbol	Meaning
N	Number of UAVs (agents)
o_i^t	Local observation of agent i at time t
h_i^t	Latent embedding of o_i^t
A^t / A_{ij}^t	Feasible radius topology (binary mask)
q_{ij}^t	Communication quality score
g_{ij}^t	Soft communication gate in $[0, 1]$ (not binary)
K / K_i	(Per-agent) communication budget / Top- K width
A_t^{AC}	Budgeted adjacency $A^t \odot E_t^K$
e_{ij}^t	Effective edge weight for attention
$G_t = (V, E_t)$	Communication graph at time t
α_{ij}^t	Attention weight on edge (i, j)
C_t / C	Instantaneous / time-averaged soft communication cost
λ_c	Communication cost weight in the joint objective
Φ_i^t	Residual guidance feature
$\beta(t)$	Residual mixing schedule
a_i^t	Action: $\pi(o_i) + \beta\Delta(\Phi_i)$

AC-DSGF Framework (three layers)



$$\text{Joint objective: } \max_{\theta} \mathbb{E}[\sum_t r_t] - \lambda_c \mathbb{E}[\sum_t C_t]$$

Fig. 2. Framework of AC-DSGF: adaptive communication (gate + budget) defines the topology consumed by residual DSGF guidance, rather than post-hoc pruning.

A. Joint Communication–Control Objective

Let $r_t = \sum_i r_i^t$ denote the team reward at time t and

$$C_t = \sum_{i \neq j} g_{ij}^t \quad (5)$$

the soft communication mass. AC-DSGF optimizes

$$\max_{\theta} \mathbb{E} \left[\sum_{t=0}^H r_t \right] - \lambda_c \mathbb{E} \left[\sum_{t=0}^H C_t \right], \quad (6)$$

i.e., communication is optimized *together with* policy learning. This differs from training a dense communicator and then pruning messages at deployment time.

B. Adaptive Communication Graph

a) *Feasible mask and quality.*: The radius mask $A_{ij}^t = \mathbb{I}[\|p_i^t - p_j^t\| < R_c]$ and quality $q_{ij}^t \in (0, 1)$ yield support $\tilde{A}_{ij}^t = A_{ij}^t q_{ij}^t$.

b) *Gate (learned topology).*: Let $h_i^t = \psi(o_i^t)$ and let ρ_{ij}^t encode an interaction-risk feature (relative closing speed / proximity prior; implemented via the edge-state channels alongside d_{ij}^t and q_{ij}^t). The communication gate is

$$g_{ij}^t = \sigma(W_g[h_i^t; h_j^t; d_{ij}^t; \rho_{ij}^t]) \cdot A_{ij}^t \in [0, 1], \quad (7)$$

where W_g denotes the parameters of a lightweight MLP (bias omitted). Unlike random dropout, g_{ij}^t is a *learned* function of agent states and risk—empirically aligned with proximity risk (Sec. VI-E).

c) *Hard budget $\Rightarrow$ AC adjacency.*: Given budget ratio $\rho \in (0, 1]$,

$$K_i = \lceil \rho |\mathcal{N}_i| \rceil, \quad E_t^K = \text{TopK}_{K_i}(\{g_{ij}^t\}_{j \in \mathcal{N}_i}), \quad (8)$$

and the adaptive communication adjacency

$$A_t^{AC} = A^t \odot E_t^K, \quad e_{ij}^t = A_{ij}^{AC} q_{ij}^t g_{ij}^t. \quad (9)$$

Training may use soft g with the cost in Eq. (6); evaluation can enforce hard Top- K without retraining (budget sweep).

d) *Difference from pruning.*: Message pruning typically removes links after (or independently of) policy learning. AC-DSGF changes the *information topology* inside the forward pass: setting $g=A$ (“no-budget”) restores near-full local broadcast and inflates C_t by over three orders of magnitude with only a small Success change (Sec. VI-D). Hence AC-DSGF is not “DSGF + ℓ_2 regularization,” but a gated topology learner under a joint objective.

C. Residual Spatial–Temporal Guidance

Sparse attention aggregates neighbors under e^t :

$$\alpha_{ij}^t \propto \exp(f(h_i^t, h_j^t)) e_{ij}^t, \quad h_i'^t = \sigma \left(\sum_j \alpha_{ij}^t W h_j^t \right), \quad (10)$$

followed by $z_i^t = \text{GRU}(h_i'^t, z_i^{t-1})$ and guidance Φ_i^t . Actions use residual decoupling

$$a_i^t \sim \pi_{\theta}(\cdot | o_i^t) + \beta(t) \Delta(\Phi_i^t), \quad (11)$$

so coordination *assists* rather than dominates execution (Table VI: 9.27% $\rightarrow$ 0.22% without residual).

D. Difference from Conventional Communication Learning

Table II contrasts AC-DSGF with common alternatives. Key distinction: the communication topology is optimized as a **decision variable** rather than treated as a fixed network structure. Joint optimization (Eq. (6)) and the explicit map $A \rightarrow g \rightarrow A^{AC} \rightarrow e$ make gating part of the forward computation, not an ℓ_2 -style afterthought on an otherwise dense GAT/DSGF stack.

TABLE II
HOW COMMUNICATION IS TREATED DURING LEARNING.

Method	Communication mechanism
Dropout / random pruning	Stochastic deletion of messages
Post-hoc / penalty-only	Constraint applied after (or around) a fixed policy graph
GAT / radius graphs	Attention on a <i>fixed</i> geometric adjacency
AC-DSGF	Task-driven gate + budget; topology is co-learned




Fig. 3. Gate stability: mean $\pm$ std communication mass C_t and mean $|\Delta E_t|$ over episodes. Sparse AC graphs are dynamic yet non-chattering.

E. Gate Stability (non-random switching)

A potential concern is that sparse gates might chatter and destabilize training. Using frozen AC-DSGF rollouts, we measure soft cost variance $\text{Var}(C_t)$ and the edge-change rate

$$\Delta E_t = ||E_t| - |E_{t-1}||. \quad (12)$$

Empirically (Fig. 3; 12 episodes), $\text{Var}(C_t) \approx 3 \times 10^{-3}$ and $>99\%$ of steps have $\Delta E_t = 0$: C_t varies smoothly with the task phase while edges rarely chatter. Communication adapts when proximity risk changes (Sec. VI-E), rather than randomly toggling.

F. Optimization

Actors use clipped PPO + GAE with the AC cost term for $\lambda_c = 10^{-3}$ (frozen). Algorithm 1 makes the Gate $\rightarrow$ Budget $\rightarrow$ Policy order explicit. Computational analysis is given in Sec. IV-G.

G. Computational and Communication Complexity

Let N be the number of UAVs, d the hidden width, $\mathcal{N}_i$ the feasible neighbors of agent i under radius R_c , and $k = \frac{1}{N} \sum_i |\mathcal{N}_i|$ the average degree ($k \ll N$ in sparse swarms). Let $K_i = \lceil \rho |\mathcal{N}_i| \rceil$ denote the Top- K budgeted degree and $K = \frac{1}{N} \sum_i K_i$ the average budgeted degree ($K \leq k$).

a) *Dense / GAT baselines.*: Full pairwise messaging or dense attention costs $\mathcal{O}(N^2d)$ per layer. Even with a binary radius mask, GAT-style aggregation that treats all in-radius neighbors equally still processes $\Theta(Nk)$ edges; under moderate R_c this rapidly approaches quadratic growth in N .

b) *AC-DSGF.*: Gate evaluation runs on the feasible edge set: $\mathcal{O}(|E|d) = \mathcal{O}(Nkd)$. Top- K budgeting per agent costs $\mathcal{O}(k \log k)$ (or $\mathcal{O}(K \log k)$ with partial selection), totaling $\mathcal{O}(Nk \log k)$. Sparse attention then operates on the budgeted support of size $\mathcal{O}(NK)$:

$$\mathcal{O}(NKd + Nd) \quad (\text{attention} + \text{GRU}), \quad (13)$$

Algorithm 1 AC-DSGF Training

Require: Observations $\{o_i^t\}$, radius R_c , budget ratio ρ (or K), weight λ_c
Require: Policy parameters θ , gate parameters ϕ (collectively θ in text)
Ensure: Updated θ, ϕ
for each training iteration **do**
 2: Collect rollouts with decentralized actors
 3: **for** each timestep t in batch **do**
 4: Embed $h_i^t \leftarrow \psi(o_i^t)$; build candidate mask A^t and quality q^t
 5: **Gate:** $g_{ij}^t \leftarrow \sigma(W_g[h_i; h_j; d_{ij}; \rho_{ij}]) \cdot A_{ij}^t$
 6: **Budget:** $E_t^K \leftarrow \text{TopK}_{K_i}(\{g_{ij}^t\})$; $A_t^{\text{AC}} \leftarrow A^t \odot E_t^K$
 7: Build $e_{ij}^t \leftarrow A_{ij}^{\text{AC}} q_{ij}^t g_{ij}^t$; $C_t \leftarrow \sum_{ij} g_{ij}^t$
 8: Aggregate neighbors with sparse attention on e^t ; $\text{GRU} \rightarrow \Phi_i^t$
 9: **Policy:** sample $a_i^t \sim \pi_\theta(\cdot | o_i^t) + \beta(t)\Delta(\Phi_i^t)$
 10: **end for**
 11: Update θ, ϕ by PPO/GAE loss plus $\lambda_c \mathbb{E}[\sum_t C_t]$
 12: **end for**

TABLE III
COMPLEXITY AND TYPICAL COMMUNICATION INTENSITY (CONCEPTUAL). K IS THE AVERAGE BUDGETED DEGREE; k IS THE AVERAGE RADIUS DEGREE.

Method	Communication	Aggregation	Complexity
MAPPO	none	—	$\mathcal{O}(Nd)$
GAT / Transformer	high	dense / full	$\mathcal{O}(N^2d)$
DSGF	medium (all radius links)	sparse k	$\mathcal{O}(Nkd + Nd)$
AC-DSGF	low (Top- K gates)	sparse $K \leq k$	$\mathcal{O}(NKd + Nd)$

which approaches linear scaling in N when K is held approximately constant by the budget ratio. Because $K \ll N$ under tight budgets, the relative advantage versus $\mathcal{O}(N^2d)$ dense attention **widens as team size grows**. Communication volume itself is bounded by $\sum_i K_i = \mathcal{O}(NK)$, whereas GAT/DSGF without gating use $\mathcal{O}(Nk)$ open geometric edges. Pairwise distance masking remains $\mathcal{O}(N^2)$ and is shared by all geometric graph methods; the practical gain of AC-DSGF is in *attended* edges and transmitted soft mass under $K \ll N$. Fig. 6 illustrates communication versus team size. Table IV reports CPU wall-clock per decision step at $N=16$ (gate overhead is small relative to learning benefits of sparse messaging).

V. EXPERIMENTAL SETUP

A. Environment and Protocol

All formal results use the VMAS continuous multi-agent navigation scenario. Unless stated otherwise:

- Training: 102,400 frames; evaluation: 200 episodes.
- Main table seeds: $\{42, 3407, 2026, 1234, 8888\}$ ($N=16$).
- Horizon $H=128$; success threshold 0.3 world units; $R_c=0.5$.
- DSGF v2 backbone and AC hyperparameters are frozen (dsgf_v2_frozen.yaml, ac_dsgf_16uav.yaml).

TABLE IV

PER-STEP INFERENCE COST ON CPU ($N=16$, 50 TIMED STEPS AFTER WARMUP). AC-DSGF ADDS A LIGHTWEIGHT GATE; SOFT EDGE MASS REMAINS FAR BELOW DENSE GUIDES.

Method	Inference (ms/step)	Soft edges (last)
GAT	1.66	—
Transformer	1.69	—
DSGF	3.35	—
AC-DSGF	3.94	≈ 0.006

TABLE V

TASK PERFORMANCE UNDER COMMUNICATION CONSTRAINTS ($N=16$, 5 SEEDS, 200-EP EVAL). AC-DSGF MAINTAINS COMPARABLE TASK COMPLETION WITH SUBSTANTIALLY REDUCED COMMUNICATION OVERHEAD (HIGHER CEI). MEAN $\pm$ STD.

Method	Success (%)	Collision (%)	Comm	CEI
MAPPO	2.40 ± 0.90	32.3 ± 1.9	0.00 ± 0.00	—
GAT	2.04 ± 0.90	29.6 ± 1.0	38.80 ± 1.99	0.0005
Transformer	0.70 ± 0.16	30.8 ± 3.0	240.00 ± 0.00	≈ 0
DSGF	4.01 ± 1.68	31.5 ± 3.3	39.27 ± 1.92	0.0010
AC-DSGF	3.95 ± 0.83	30.8 ± 1.4	0.43 ± 0.09	0.0913

B. Baselines

MAPPO (no graph messaging); GAT-MAPPO (binary radius graph + alignment reward); Transformer / full attention; DSGF v2 (quality-weighted sparse graph + residual); **AC-DSGF** (DSGF v2 + learned gate + $\lambda_c C_{\text{comm}}$).

C. Metrics

Success S , collision rate, episodic reward, communication cost C_{comm} (Eq. (3)), and CEI = $S/(C_{\text{comm}} + \varepsilon)$ (Eq. (4)). Demo videos are qualitative and *not* used as Table metrics.

D. Experiment Map

- 1) **Table V**: task performance under communication constraints ($N=16$).
- 2) **Table VI**: residual ablation ($N=4$).
- 3) **Table VII** / **Fig. 4**: communication budget analysis.
- 4) **Table VIII**: silence-collapse validation.
- 5) **Fig. 5** / **Fig. 3**: behavior & gate stability.
- 6) **Fig. 6** / **Fig. 8**: scalability & Pareto.

VI. RESULTS

A. Task Performance under Communication Constraints

AC-DSGF achieves Success *comparable* to DSGF (3.95% vs. 4.01%) while reducing communication from 39.27 to 0.43 ($\sim 91\times$) and lifting CEI by nearly two orders of magnitude (Fig. 8). We deliberately **do not** claim Success leadership: the value proposition is communication-efficient cooperation under a realistic swarm constraint.

B. Ablation Study

This ablation is not merely “remove a module $\Rightarrow$ score drops.” It isolates a failure mode of guided MARL: when Φ dominates the actor, policies overfit coordination surrogates and lose long-horizon task competence. Residual decoupling keeps guidance *assistive*, explaining the $\sim 42\times$ Success collapse without it. (AC gate ablations appear in Sec. VI-D.)

TABLE VI

DSGF RESIDUAL ABLATION AT $N=4$ (MECHANISM TEST; *not* TABLE I SCALE).

Variant	Success (%)
w/o Dynamic Graph	0.28
w/o Temporal	0.86
w/o Residual	0.22
Full DSGF	9.27

TABLE VII

SUCCESS (%) UNDER COMMUNICATION BUDGET RATIO ρ ($N=16$, SEED 42, 64-EP DIAGNOSTIC EVAL; QUALITATIVE CONSISTENCY WITH TABLE I CEI RANKING). SOFT-COST SCALES DIFFER FROM TABLE I; WE REPORT TRENDS, NOT ABSOLUTE % INTO THE MAIN CLAIM.

Budget ρ	GAT	DSGF	AC-DSGF
100%	21.3	21.7	27.1
50%	18.2	21.6	25.2
10%	16.9	21.7	24.2

C. Communication Budget Analysis

Under a 90% budget cut (100% $\rightarrow$ 10%), AC-DSGF success drops only $\sim 11\%$ relatively (27.1 $\rightarrow$ 24.2), i.e., *graceful degradation*. Paper wording prefers “over three orders of magnitude” for gate-level soft-cost reductions versus “no-budget” ablation, rather than advertising un-normalized $5000\times$ factors without defining soft cost.

D. Silence Collapse Ablation

Interpretation. Sparsity is *not* silence collapse. AC-full trades a few percentage points of Success for over-three-orders reduction in Comm relative to full local broadcast—matching the locked claim of comparable performance at far lower messaging.

E. Communication Behavior Interpretation

Pearson correlations with Comm: goal distance 0.03, dispersion -0.57 , nearest-neighbor risk **0.84**. Communication intensifies under proximity risk and attenuates when agents disperse—evidence of *task-dependent* topology rather than random pruning.

F. Scalability Analysis

Fig. 6 places the Table I communication result of AC-DSGF in a scaling context: maintaining task effectiveness does not require paying the dense-graph communication tax as N grows.

G. Robustness under Packet Loss

From $p=0$ to $p=0.7$, GAT Success drops from 21.2% to 15.8% (relative $\sim 25\%$), while AC-DSGF moves only from 26.2% to 25.8%. MAPPO is insensitive by construction (no graph messaging). We interpret AC-DSGF’s flat curve as **graceful degradation**: messaging was already sparse and task-conditioned, so random drops remove fewer coordination-critical edges on average.




Fig. 4. Success vs. communication budget ($N=16$). AC-DSGF maintains task effectiveness under severe budget cuts (graceful degradation).

TABLE VIII

COMMUNICATION MECHANISM VALIDATION ($N=16$, AC CHECKPOINT). OPENING ALL RADIUS EDGES INFLATES COMM MASSIVELY WITH ONLY A SMALL SUCCESS GAIN.

Variant	Success (%)	Comm	CEI
AC-full (learned)	24.6	0.008	32.0
AC-no budget ($g=A$)	28.9	40.6	0.007
AC-random gates	28.5	20.3	0.014

H. Communication–Performance Pareto

LIMITATIONS

Absolute Success at $N=16$ remains low for all methods; claims are comparative and communication-centric. Budget/silence diagnostics may use different episode counts than Table I and must be described as such. Soft gate costs are not identical to packet counts on hardware radios.

VII. CONCLUSION

We presented AC-DSGF for communication-constrained multi-UAV cooperative navigation. By jointly optimizing task return and communication mass $\mathbb{E}[\sum r_t] - \lambda_c \mathbb{E}[\sum C_t]$, learning a state-/risk-dependent gate, and enforcing Top- K budgets atop residual DSGF guidance, AC-DSGF **maintains task effectiveness while reducing unnecessary communication**. At $N=16$, it achieves Success comparable to DSGF with approximately $90\times$ lower communication cost, with supporting budget, silence-collapse, trigger, compute-cost, and packet-loss analyses. Future work includes hardware radio models and larger-scale field validation.

REFERENCES

- [1] C. Yu, A. Velu, E. Vinitsky, J. Gao, Y. Wang, A. Bayen, and Y. Wu, “The surprising effectiveness of PPO in cooperative multi-agent games,” in *NeurIPS*, 2022.
- [2] R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch, “Multi-agent actor-critic for mixed cooperative-competitive environments,” in *NeurIPS*, 2017.
- [3] T. Rashid, M. Samvelyan, C. Schroeder, G. Farquhar, J. Foerster, and S. Whiteson, “QMIX: Monotonic value function factorisation for deep multi-agent reinforcement learning,” in *ICML*, 2018.
- [4] S.-J. Chung, A. A. Paranjape, P. Dames, S. Shen, and V. Kumar, “A survey on aerial swarm robotics,” *IEEE Trans. Robot.*, vol. 34, no. 4, pp. 837–855, 2018.
- [5] M. Bettini, R. Kortvelesy, J. Blumenkamp, and A. Prorok, “VMAS: A vectorized multi-agent simulator for collective robot learning,” *arXiv preprint arXiv:2207.03530*, 2022.
- [6] J. Jiang and Z. Lu, “Learning attentional communication for multi-agent cooperation,” in *NeurIPS*, 2018.




Fig. 5. Trigger analysis: Comm correlates strongly with proximity risk ($\text{corr}=0.84$) and negatively with spatial dispersion ($\text{corr}=-0.57$).




Fig. 6. Communication versus number of UAVs. Dense / GAT schemas grow quickly; DSGF uses geometric support; AC-DSGF at $N=16$ sits far lower (Table V), consistent with $O(NK)$ budgeted messaging (Sec. IV-G).

- [7] S. Iqbal and F. Sha, “Actor-attention-critic for multi-agent reinforcement learning,” in *ICML*, 2019.
- [8] J. Jiang, C. Dun, T. Huang, and Z. Lu, “Graph convolutional reinforcement learning,” in *ICLR*, 2020.
- [9] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, “Graph attention networks,” in *ICLR*, 2018.
- [10] A. Singh, T. Jain, and S. Sukhbaatar, “Learning when to communicate at scale in multiagent cooperative and competitive tasks,” in *ICLR*, 2019.
- [11] R. Wang, X. He, R. Yu, W. Qiu, B. An, and Z. Rabinovich, “Learning efficient multi-agent communication: An information bottleneck approach,” in *ICML*, 2020.
- [12] J. Foerster, Y. M. Assael, N. de Freitas, and S. Whiteson, “Learning to communicate with deep multi-agent reinforcement learning,” in *NeurIPS*, 2016.
- [13] S. Sukhbaatar, A. Szlam, and R. Fergus, “Learning multiagent communication with backpropagation,” in *NeurIPS*, 2016.
- [14] A. Das, T. Greeff, A. Gupta *et al.*, “TarMAC: Targeted multi-agent communication,” in *ICML*, 2019.
- [15] D. Kim, S. Moon, D. Hostallero *et al.*, “Learning to schedule communication in multi-agent reinforcement learning,” in *ICLR*, 2019.
- [16] I. Mordatch and P. Abbeel, “Emergence of grounded compositional language in multi-agent populations,” in *AAAI*, 2018.




Fig. 7. Task success under random edge drops ($N=16$, 64-ep diagnostic). AC-DSGF remains nearly flat as loss increases, whereas GAT degrades more noticeably—consistent with *graceful degradation* under unreliable links.

TABLE IX

SUCCESS (%) VS. PACKET-LOSS RATE p ($N=16$, SEED 42, 64 EPS). WE EMPHASIZE TREND (DEGRADATION SLOPE), NOT ABSOLUTE COMPARISON WITH TABLE I.

Method	0%	10%	30%	50%	70%
MAPPO	31.4	30.7	31.2	33.5	29.9
GAT	21.2	17.6	16.1	14.6	15.8
DSGF	21.8	24.2	20.7	21.5	18.5
AC-DSGF	26.2	26.1	26.4	28.3	25.8




Fig. 8. Communication–performance Pareto ($N=16$, Table V). AC-DSGF occupies the high-efficiency region: comparable Success, minimal Comm.